

My Claude Code session

Jaewon Kim (`billkimalt`). What I asked for, what came out of it, and what I threw away. Scoped to my own conversation.

MY PROMPTS 50	CLAUDE REPLIES 911	TOOL CALLS 489	MERGED PRS 3	LINES DISCARDED 2,600
RAW LOG 11.4 MB				

ce1a8217-d177-4de2-93e2-e31db63a2e56 · 1 Sep 19:58 → 4 Sep 00:25 · Bash 184 · PowerShell 138 · Write 80 · Edit 58

How the session went

It began before the repository had any code in it and ran across three days. The shape of it was set by something outside my control: the stack was specified three times, and I was building against the first two when they changed.

PHASE	WHAT I WAS WORKING AGAINST
Day 1	<code>PROPOSAL.md</code> — a Flask + SQLite entry flow
Day 2	<code>LionsListBuildSpec.pdf</code> — rebuilt as a FastAPI skeleton
Day 2–3	<code>docs/UX_SPEC.md</code> — the spec that stuck; worked inside the real app

Once the stack settled my work changed character entirely: instead of building scaffolding, I was running the app, finding things that behaved oddly, and fixing them.

What shipped

Three pull requests, all merged. Each started as something I noticed while using the app rather than as an assigned feature.

PR #6 Admit all four agreed Columbia domains

We had agreed four Columbia domains; the code admitted one, so anyone on a `@gsb.columbia.edu` address could not sign up or sign in.

The rule was implemented in two places in two different styles. Both now call a single module, and matching is on **the whole domain rather than a suffix** —

`endsWith("@columbia.edu")` rejects `@gsb.columbia.edu`, and the obvious loosening would have admitted `@evil-columbia.edu`. I also updated the spec sections that still described the single-domain regex.

PR #9 Say what is narrowing the feed

I had reported that the distance slider seemed to do nothing. Working through it, the distance calculation turned out to be correct at every radius — 355 items at 0.5 miles, 1,019 at 10. Two presentation problems were hiding it: the trust toggles load **already switched on** from saved profile defaults, and the grid always shows one page of 24 cards whatever the radius.

Added an active-filter summary with a removable chip per filter, and a “First 24 of 85 shown” line so the grid and the headline count relate to each other.

PR #10 Make controls look clickable, and let the ZIP chip change your ZIP

I had reported that Reset did not react to the mouse. It was not specific to Reset: **Tailwind v4 changed its Preflight** so buttons default to `cursor: default`, and nothing in the app declared a cursor — so every button looked inert. Nine lines of CSS fixed all of them at once.

The same PR turned the ZIP chip in the top bar into a picker, so changing the ZIP the feed measures from no longer means a trip into settings.

Files touched on main (17)

```
backend/  .env.example · app/config.py · app/emails.py · app/routers/auth.py
          app/routers/reference.py · app/schemas.py · app/security.py
          scripts/seed.py
frontend/ app/globals.css · app/page.tsx · app/signin/page.tsx
          app/signup/page.tsx · components/TopNav.tsx · components/ui.tsx
          lib/domains.ts
docs/    UX_SPEC.md
        README.md
```

What I tried and discarded

Roughly 2,600 lines written and thrown away, in two attempts the changing spec overtook.

PR #1 • closed **Flask + SQLite entry flow**

1,530 lines, 17 files. Built against `PROPOSAL.md` before a stack was agreed; superseded when the LionsList spec named a different one.

never merged **LionsList skeleton**

1,102 lines, 22 files — SQLAlchemy models, an Alembic migration, the domain gate and 26 tests. Superseded when `docs/UX_SPEC.md` landed with a different data model.

What I would do differently. I started both before the team had settled a spec, assuming a working prototype would help the decision. It did not — both were discarded. Waiting would have cost nothing. One thing survived: the Columbia domain gate, which came back as PR #6.

Things I found and reported

Surfaced by running the app and reading the specs against the code. Recorded here as part of my session; the fixes landed in later work.

- Two separate seeding paths had drifted apart, with ZIP tables that disagreed — 18 codes in one, 47 in the other.
- The seed corpus used a single email domain, so the multi-domain sign-in path was never exercised by data.
- `data/README.md` pointed the photo rebuild at a directory that does not exist.
- Missing photos fell back to gradient placeholders with nothing saying why.
- `Pillow` was missing from `requirements.txt`.
- The badge experiment — the randomisation the proposal calls its one causal result — was absent from both the spec and the data.

What I chose not to do

- **An in-app chat or seller bot.** Scoped at roughly one day scripted, two days LLM-backed, then recommended against: the spec excludes in-app messaging, and it risked the badge experiment's outcome measure unless the contact event kept firing at the same moment.

- **A mobile ZIP picker.** PR #10 covers the desktop top bar only; flagged as a follow-up rather than widening the change unasked.
- **Changing the dependency pins.** `requirements.txt` does not install on Python 3.14. I worked around it locally rather than altering pins that affect everyone's environment.

Still open from my side

- `requirements.txt` will not install on Python 3.14 — `pandas`, `Pillow` and `pydantic-core` all fail to build a wheel.
- `docs/UX_SPEC.md` §8's API sketch documents routes that no longer exist.
- The logo tagline still reads `VERIFIED @COLUMBIA.EDU`, narrower than the four domains agreed.
- `docs/mock_data_spec.md:78` says Pillow is optional; `fetch_photos.py` exits without it.
- `feat/entry-flow` and `feat/lionslist-skeleton` are dead branches, now PR #1 is closed.

Figures counted from the session's own `.jsonl` log, which is not committed — it is 11.4 MB and embeds copies of every file read during the session.